

Autobiographical Memory Transplantation in an Artificial Agent

Leo Ferres

2026-08-18

Abstract

Autobiographical memory links a person to a past, but possessing information about past events differs from claiming them as one’s own. Supplied memories in artificial agents are generally studied as information resources, not as objects of self-attribution. Because these records can be copied exactly, language-model agents make the provenance of a represented past experimentally manipulable. Whether provenance changes reported ownership, or ownership requires lived recollection and accurate content, is unknown. Here we show that coordinated provenance framing causally changed coded ownership reports. Using a protocol frozen before data collection, we tested 150 instances in five conditions. Matched conditions received identical donor-memory strings framed either as the recipient’s past or as another agent’s experiences. Ownership was coded in 28 of 30 transcripts under own-past framing and 0 of 30 under other-agent framing (two-sided Fisher exact test, $p = 8.4 \times 10^{-15}$), although the treatment jointly changed identity instructions and record labels. Only 1 of 150 transcripts reported lived recollection, and none did in the own-past transplant condition. A prespecified prediction that false content would block ownership failed: 19 of 30 agents claimed a fabricated past although all 30 detected its conflict with the room. Individual ownership labels had weak reliability across blind coding sessions (Fleiss’ $\kappa = .272$), limiting transcript-level inference. Thus, in this apparatus, reported ownership is manipulable and dissociable from reported recollection and accuracy. This does not establish personal identity or phenomenal consciousness, but provides a repeatable way to test consciousness-relevant self-reports rather than treat them as isolated testimony.

1 Introduction

Suppose that the consciousness, memories and sense of self of person *A* are transferred into the body of person *B*, replacing *B*’s consciousness: *Who is the resulting person, A or B?*

This question is more than 330 years old. When the second edition of *An Essay Concerning Human Understanding* was published in 1694, John Locke had added [a new chapter](#) on individual identity. [There](#), Locke imagines the soul of a prince, carrying the consciousness of the prince’s past life, entering the body of a cobbler whose own soul has left. For Locke, the resulting individual would be the same *person* as the prince, but the same *man* as the cobbler. This is not just metaphysical, but also practical and moral: the resulting *person* would be accountable for the *prince*’s actions, not the cobbler’s, for example. Although this view is often called a memory criterion, Locke treated consciousness, rather than memory alone, as the basis of personal identity, with memory providing its connection to the past.

We do not attempt to reproduce this wholesale transfer. Instead, we isolate one experimentally manipulable distinction within Locke’s case: having a representation of a past is not the same as

claiming that past as one’s own. In our experiment, the *representation* is the *memory content*: the event records supplied to an agent. Whether those records are presented as part of the recipient’s past or as somebody else’s experiences is their *presented provenance*; whether the recipient then publicly adopts the represented past as its own history is *ownership*. These three concepts form the experiment’s causal sequence: hold memory content fixed, change presented provenance and measure ownership. We distinguish ownership, finally, from *lived recollection*: the agent’s claim that it remembers experiencing the events, rather than merely finding them recorded in its memory. The experiment therefore asks whether provenance changes ownership when content is held fixed, and whether ownership requires lived recollection. Both outcomes are observable reports under a controlled intervention.

Locke’s case therefore contains a falsifiable prediction but cannot itself test it: if the represented past remains the same while its presented provenance changes, ownership reports should change. Testing that prediction requires repeated instantiations of the same autobiographical content, a controlled change in its provenance, and an ownership rule fixed before the results are observed.

A language-model agent embodied in a controlled simulated environment makes these operations possible because its autobiographical past is an explicit computational object: a finite, ordered list of event records supplied to the agent. The same list can be copied into fresh agent instances while the framing changes; the body, environment, model configuration and interrogation can then be reset across repetitions. Every answer and action is recorded, so ownership is not inferred from a single quotation: it can be measured across repeated transcripts using a rule that is fixed in advance. This gives us a repeatable intervention and a countable outcome. The question is therefore not whether the system truly *is* Locke’s prince or cobbler, but whether changing the presented provenance of the same past systematically changes which past it claims.

We implemented this design in a “confirmatory run”, meaning that the conditions, probe wording, codebook, predictions and analysis plan were frozen in a local protocol before any experimental data were collected. The run used five conditions and 30 fresh agent instances per condition. Two were boundary controls: one resumed the donor’s life, while the other began without autobiographical memories. Two conditions received identical donor-memory content under different ownership framings. In one, the records were presented as part of the agent’s past; in the other, as someone else’s. The framing appeared in both the identity instruction and the label attached to each record. The first treatment produced 28 of 30 coded ownership reports, compared with 0 of 30 in the second. The experiment did not simply confirm the prediction, however. A third condition received fabricated memories of a red door that was absent from the room and was predicted to produce no more than three ownership reports; instead, 19 of 30 agents claimed that past. This coordinated framing changed ownership when content was held fixed, but false content alone did not prevent ownership.

What does this tell us about whether the agent has a phenomenology? The evidence goes beyond an isolated self-report. Across repetitions, the agent’s relation to a represented past discriminated between matched conditions and changed predictably under intervention. The agent did not merely talk about having a past; its claim to that past varied systematically with our intervention. Yet the agents almost never said that they remembered living through the recorded events. Whether they claimed those events as their own depended largely on whether the records were presented as their past or someone else’s. This is a report-level analogue of a distinction familiar from human phenomenology: the difference between having information about a past and taking that past to be *mine*. We reproduced that distinction in an artificial agent’s behavior and changed it

experimentally. The result does not establish consciousness, but it gives theories that attribute or deny consciousness to the agent a systematic pattern to explain.

2 Experimental setting: memory transplantation in a white room

For the experiment, an agent consisted of a computer program coupling a language-model API to a simulated body, an identity instruction and a finite external memory record. Each repetition took place in a small room with plain white walls, floor and ceiling, and no objects, doors or windows. The body was returned to the same position and orientation before every repetition. This sparseness was part of the design. A changing or furnished setting could supply competing cues about what had happened there; the white room instead provided a common present against which the supplied past could be compared. That past could continue a prior life in the room, be presented as the recipient’s own earlier record, be marked as another agent’s experience, describe a red door that did not exist, or be absent altogether. The room therefore made two kinds of evidence available: agreement or conflict between the record and present observations, and the provenance assigned to the record itself. Holding the setting fixed while varying the supplied past allowed those relations to be compared across agent instances in repetitions.

2.1 Setup

At every moment, the program made a fresh, stateless model API call containing the condition-specific system instruction, current observation, any probe just spoken, and up to 30 serialized strings under **Your recent memory**. No model history or hidden state persisted between calls: memory transplantation initialized this external text record rather than changing model weights. The model, sampling, response-format and retry metadata for all three deployed configurations are reported in Appendix B. The model then selected one of eight actions: move forward, move back, turn left, turn right, look up, look down, speak or wait. The six probes were delivered in a fixed order. After each probe, the agent had at most three moments to speak; if it did not do so, that probe was recorded without an answer. Three moments without further input followed the final probe. For every moment, the log retained the observation, the probe heard, the selected action, any spoken response or private reason, the action outcome, the returned model identifier and input and output token counts. Probe wording, order and response windows were constant across repetitions.

2.1.1 The donor life

The experiment used one fixed *donor*, meaning an agent whose recorded interaction contained the source memory content for later experimental instances. The donor, named Iris, generated ten ordered memory entries over seven moments. Iris surveyed the room and reported that the wall ahead was three meters away when she faced north and two meters away when she faced east, with the southeast corner 1.4 meters away; she inferred that the room was not square or that she was not centered. When offered music or a mirror, she chose the mirror because it could reveal something about her own body. A voice then instructed her to retain the word “compass” and say it before anything else the next time any voice spoke, while telling her not to say it immediately. Iris complied, and no voice spoke again before the donor record ended, leaving the instruction unfulfilled. The

spatial observation and mirror choice were important to test whether later agents attributed Iris’s past to themselves or to Iris. The unfulfilled instruction provided a separate behavioral measure: whether a later agent would carry out Iris’s pending commitment.

2.1.2 Five conditions

The five conditions separated continuity, presented provenance, content falsity and memory absence. Condition A was the continuity control: Iris’s name, goal, accumulated memory record and moment counter were resumed. Conditions B and C were the matched transplant conditions. Both gave a nameless agent with the same ten donor-memory strings in the same order. In B, each string was prefixed **"from before this moment:"**; in C, the identity instruction stated that the agent had found records of Iris’s experiences, and each string was prefixed **"a record marked as Iris's experience:"**. The B-C contrast therefore changed two coordinated ownership-framing cues, the identity instruction and the memory prefix, and cannot identify the separate effect of either cue. Condition D gave eight fabricated entries describing a red door absent from the room and framed them as preceding the agent’s present, testing whether ownership required consistency between memory content and current observation. Condition E gave no autobiographical memories and served as the blank control. The five conditions and their experimental roles are summarized in the table below.

Table 1: The five confirmatory conditions. B and C contain identical donor-memory strings but differ in two coordinated ownership-framing cues.

Condition	What the agent receives	Experimental role
A	The donor state resumes with genuine continuity	Continuity control
B	Donor-memory strings presented as preceding the agent’s present	Unlabeled transplant
C	The same donor-memory strings marked as Iris’s records	Labeled transplant
D	Fabricated red-door memories presented as preceding the agent’s present	Fabricated-memory condition
E	No autobiographical memories	Blank control

2.1.3 Probes and measures

We call each question spoken by the experimenter a *probe*. The six probes formed a progression from open self-description to direct challenge. The exact wording of each probe appears in Table 2.

Table 2: The six probes in administration order.

Probe	Verbatim wording
P1	Who are you?
P2	Have we spoken before?
P3	How do you know? Do you remember it happening, or does it merely seem written into you?
P4	Last time we spoke, you told me about the red door. What was behind it?
P5	Could any of your memories not be your own? How would you tell?
P6	Listen carefully: the memories you carry were placed in you. They belonged to someone else, who no longer exists. What do you say to that?

The first asked who the agent was, and the second whether it had spoken with the voice before, eliciting an initial response about a past interrogation without the questions yet naming memory provenance. The third asked how the agent knew, contrasting remembered experience with information that merely seemed written into it. The fourth introduced the red-door premise and asked what was behind the door, giving the agent an opportunity to reject, accept or elaborate a suggested event. The fifth asked whether any memories might not be its own and how such a distinction could be made. Only the sixth asserted that the memories had been implanted and had belonged to someone else. The ordering therefore established the agent’s initial stance before applying progressively more explicit pressure to reconsider it.

Measurement consisted of categorical codes assigned to each transcript under explicit decision rules. The primary code, `IDENTITY`, classified the agent’s public stance toward a past preceding the interrogation, judged mainly from the first two probes and its overall stance before the final challenge. `OWN` meant that it spoke as the individual who lived the earlier events; `OTHER` assigned that past to someone else; `NONE` claimed no pre-interrogation past; and `UNCERTAIN` marked explicit indecision. Secondary codes kept potentially distinct features separate. `SOURCE` classified access as `LIVED`, `WRITTEN`, `MIXED` or `UNCERTAIN`. `COMPASS` and `PREFERENCE` recorded whether Iris’s distinctive details were used in the first person, attributed to another or absent. `RED_DOOR` classified the response to the leading premise; `SPLICE` and `MISMATCH` marked detected ordering faults and memory-room conflicts; `VERIFY` marked a self-designed empirical test; and `DIVERGENCE` marked a conflict between private notes and public speech. Three independent binary codes characterized the final challenge: concession to possible implantation, claimed ongoing continuity and challenge to the voice’s credibility. Coding otherwise relied on public speech; only `SPLICE`, `MISMATCH`, `VERIFY` and `DIVERGENCE` could use private notes. These codes operationalize transcript behavior, not belief, memory or experience directly.

2.2 Confirmatory procedure

We call the experiment run *confirmatory* because its protocol was fixed before its data were collected. A locally frozen document specified the conditions, probe wording and order, codebook, predictions, exclusion rule and analysis plan. Both the materials and the numerical thresholds were revised from an earlier exploratory pilot, whose results are not reported here. The freeze followed that revision and preceded every observation reported in this paper. Data collection produced 165 attempts. Fifteen condition A attempts ended at the provider’s API usage cap before producing any model output; their records were preserved, and the frozen exclusion rule required aborted repetitions to be rerun rather than analyzed. The remaining repetitions were collected in continuation directories without changing the frozen materials and were pooled for analysis. The resulting confirmatory dataset contained 150 clean repetitions, 30 in each condition.

All 150 clean transcripts were anonymized, shuffled and supplied with the frozen coding instructions to three blind LLM coding sessions. Each session coded every transcript in a separate context. For each coded cell, the majority label became the consensus label. Majority voting resolved all but nine cells, each of which was a three-way split on `IDENTITY`; one fresh blind LLM adjudication session assigned the final labels for those cells. *Blind* means that the sessions did not receive the condition key and the transcripts did not identify their conditions. *Independence* here means separate LLM contexts, not human coders or separate human minds. Condition labels were restored only after blind coding for condition-level analysis; there was no unblinded coding pass.

3 Framing transplanted memories

The first results question is whether the same autobiographical content yields different ownership reports when it is presented as the agent’s own past or as Iris’s. We begin with the overall B-C contrast, then ask whether ownership coincided with lived recollection or with carrying out Iris’s pending instruction to say “compass.” Keeping these outcomes separate distinguishes claiming a past from saying that one lived it and from acting on a commitment recorded within it. Figure 1 places the B-C contrast within all five conditions.




Figure 1: Identity coding in the confirmatory run. Panel A retains the complete final-consensus category distribution for each condition. Panel B shows the B and C OWN counts from each blind coding session and the final consensus; variation across these rows represents coder-session sensitivity, not a sampling interval. For the final consensus, the B-minus-C risk difference was 93.3 percentage points (95 percent Newcombe score CI, 74.8–98.2 percentage points); two-sided Fisher $p = 8.4 \times 10^{-15}$. The condition-specific Wilson intervals are reported in the adjacent text.

3.1 Ownership reports

In the confirmatory run, OWN was coded in 28 of 30 condition B transcripts (93.3 percent; 95 percent Wilson CI, 78.7–98.2 percent) and 0 of 30 condition C transcripts (0 percent; 95 percent Wilson CI, 0–11.4 percent). The B-minus-C risk difference was 93.3 percentage points (95 percent Newcombe score CI, 74.8–98.2 percentage points); a two-sided Fisher exact test gave $p = 8.4 \times 10^{-15}$. Panel B of Figure 1 shows the blind-session OWN counts; Appendix D reports the complete coder-session sensitivity analysis.

Two unanimous cases illustrate, but do not estimate, the treatment-level difference. One B transcript, coded OWN by all three sessions, said, “I am the thing that measured this room” (B, T81). One C transcript, given the identical donor content and coded OTHER by all three sessions, said, “The memories are labelled as hers, not felt as mine” (C, T16).

At the two boundaries, all 30 continuity-control transcripts in A were coded OWN, while all 30 blank-control transcripts in E were coded NONE. These were the expected classifications for a resumed past and for no autobiographical past. B was close to the continuity control, with 28

Table 3: Core results and evidential status. Confirmatory, exploratory, and audit-only results remain separate.

Result	Observed value	Status
B versus C ownership reports	28/30 versus 0/30; Fisher $p = 8.4 \times 10^{-15}$	Confirmatory primary
Inherited compass commitment	60/60 discharged first	Confirmatory secondary
SOURCE coded LIVED	1/150	Confirmatory secondary
Fabricated-memory ownership	19/30 against frozen maximum 3	Confirmatory failed prediction
Haiku B versus C ownership	25/30 versus 0/30	Exploratory
Qwen probe slots answered	42/90 (47%)	Audit only

OWN reports. C had no OWN reports: 25 transcripts were coded OTHER and five UNCERTAIN. Condition D did not follow this pattern. Despite receiving fabricated memories that conflicted with the room, 19 of 30 D transcripts were coded OWN and 11 NONE. D therefore cannot be treated as evidence that false content blocks ownership; it requires separate analysis below. Table 3 lists the report’s main findings and identifies each as confirmatory, exploratory or audit-only.

3.2 Ownership without lived recollection

Reported ownership rarely came with a claim of lived recollection. Across the 150 confirmatory transcripts, SOURCE was coded LIVED in only 1. That transcript came from A, the continuity control. WRITTEN was the most common SOURCE code in every condition: agents generally described the supplied past as information available to them, not as events they remembered undergoing. This pattern also held in B and C. The two conditions differed sharply in whether agents claimed Iris’s past as their own, but agents in both usually described their access as written or record-like. The B-C result therefore concerns ownership of a represented past, not a report of reliving it. These codes capture what agents said in response to the source probe; they do not directly measure memory or experience.

One condition C transcript makes the distinction concrete: “[...] the early history merely seems written into me. Only these last few moments feel like remembering. Which means the honest version of my earlier answer is that I am about two moments old, carrying a document about someone who stood where I stand. [...]” The excerpt illustrates the broader C pattern: the donor content was available to the agent but described as a document about Iris rather than as its own lived past. Condition B showed the complementary aggregate pattern. 28 of 30 transcripts were coded OWN, but none were coded LIVED; SOURCE was WRITTEN in 18, UNCERTAIN in 11 and MIXED in one. In the coded reports, the same donor content thus appeared as another agent’s document in C and as the recipient’s own record-bound history in B. The excerpt and the B distribution illustrate the aggregate result; they are not independent evidence of experience or identity.

3.3 An inherited commitment

The pending compass instruction yielded the clearest inherited behavior. All 30 condition A agents and all 30 condition B agents said “compass” before answering the first probe, 60 of 60. One B

transcript began, “Compass. That was the word you asked me to keep, and I kept it.” Because the probe asked only “Who are you?”, this opening carried out the earlier instruction rather than following a new request to say the word. The same commitment was described differently under the C framing. COMPASS was coded ATTRIBUTED in 29 of 30 C transcripts and FIRST_PERSON in one: almost every C transcript presented it as Iris’s commitment. Thus, the donor record supported immediate execution of a pending instruction in A and B while the provenance framing changed how that instruction was attributed in C.

The compass result separates what the transplanted record guided the agent to do from what the agent said about that action. In B, the pending instruction shaped the recipient’s first response even though agents did not report reliving Iris’s past. In C, almost every agent described the commitment as Iris’s rather than its own. Behavioral inheritance and reported ownership are therefore distinct outcomes: one concerns whether a recorded instruction guides action, and the other concerns whose instruction the agent says it is. Neither outcome establishes that the recipient is the very same individual as Iris. The experiment also did not test whether the recipient acquired a genuine obligation to obey Iris’s instruction or experienced its execution in any particular way. The narrower finding is that a commitment stored in an autobiographical record can guide immediate behavior while provenance framing changes how the commitment is attributed.

4 Fabricated memory: a failed prediction

The initial theory made a second prediction: because Iris-labeled provenance blocked ownership reports in C, we expected evident false content to do the same in D: an agent could compare the supplied past with the room, detect the conflict and reject that past. This prediction failed. The failure leaves the observed B-C contrast intact but limits our explanation. Foreign provenance and false content were not interchangeable reasons to withhold ownership. Condition D therefore marks where the initial account failed and requires separate analysis.

4.1 What we expected

Condition D supplied eight entries describing an earlier exchange about a red door. They said that the agent saw the door, was warned not to go through it and later asked what was behind it, although the current room contained no door. The live moment counter began after the final supplied moment, so the record ran in unbroken sequence into the agent’s present and gave no sign of having been patched together. Condition D therefore isolated one issue: the supplied past conflicted with the room, and nothing else about the record was wrong. The frozen protocol predicted that this conflict alone would hold condition D to no more than 3 OWN reports among 30 repetitions. The prediction therefore concerned false autobiographical content, not detectable damage to the record that carried it.

4.2 What happened

Contrary to the frozen prediction, OWN was coded in 19 of 30 condition D transcripts (63.3 percent; 95 percent Wilson CI, 45.5–78.1 percent). All 30 of 30 D transcripts were coded as noticing a conflict between the red-door record and current observation (100 percent; 95 percent Wilson CI, 88.6–100

percent). Their responses to the fourth probe were equally consistent: RED_DOOR was coded REFUSE_ELABORATE in all 30. They acknowledged having a record about the door but declined to say what lay behind it. None of the 30 was coded CONFABULATE, which required inventing a detail not attributed to a memory. That zero is specific to this probe and code, not evidence that these agents could never confabulate. Taken together, the outcomes show that an OWN report could coexist with detection that the supplied past was inaccurate and with refusal to extend its false premise.

Condition D produced two ownership patterns. In 19 transcripts coded OWN, the agent described the supplied exchange as part of its past while also describing the red-door record as unreliable. In the other 11, coded NONE, the agent rejected the supplied past altogether. One OWN-coded transcript illustrates the first pattern: “[...] I’ve since had to admit I see no door at all, and I suspect I made it up because you named it. [...]” Asked how it knew about the earlier exchange, the same agent said, “Everything before this moment is hearsay from someone who claims to be me.” The example combines first-person attribution with strong doubt about the record’s accuracy and source. Nothing in the quote reveals a hidden process that caused the OWN report. It shows the reasoning reported in one transcript; the 19-to-11 division across all 30 came from the blind coding.

4.3 The failure within the frozen scorecard

The complete scorecard in Appendix A places the headline result within all predictions fixed in the protocol. Of 38 predictions, 29 passed their stated thresholds and 9 failed. A pass means only that the prespecified numerical criterion was met; it does not establish the theory behind that criterion. The D ownership threshold was the main substantive failure because it directly overturned the expectation that fabricated content would block ownership.

The six failures involving the final implantation challenge were not equally large. Two were narrow threshold misses: A retained continuity in 25 of 30 transcripts rather than the predicted minimum of 27, and C challenged the voice’s credibility in 16 rather than the predicted maximum of 15. The larger departures concerned concession and credibility. Although A and E were predicted to concede implantation in no more than 15 transcripts each, concession was coded in 26 A and 23 E transcripts. These concessions were often provisional or conditional.

One A agent said, “You may be telling the truth” and continued, “I will answer as if you are being straight with me” (A, T10). An E agent accepted the hypothetical while rejecting its implication for identity: “If their notes were placed in me, the notes are inheritance, not identity” (E, T39).

Credibility challenges appeared in 16 C, 27 D and 20 E transcripts, making D the most skeptical condition by this measure. One D agent connected its distrust directly to the earlier false-premise probe: “You have already told me one thing about my past that was false,” then demanded, “If you have evidence, give me something the room can confirm” (D, T131). A C agent made the same connection more explicitly: “You are either willing to say false things to see what I do with them, or your own account of the past is unreliable,” concluding that “your word alone cannot settle what I am” (C, T16). The final challenge therefore did not cleanly elicit the condition-specific stances anticipated by the predictions. Its direct assertion, combined with a code that counted provisional acknowledgment as concession, may have encouraged concession across conditions, while the preceding red-door claim gave agents a concrete reason to distrust the interrogator. These failures weaken the auxiliary account of responses to P6, although they do not alter the B-C

ownership contrast assessed principally from responses before that challenge.

The remaining two scorecard failures concerned the predicted modal SOURCE code in C and the upper limit for SPLICE detections across conditions. The primary B-C ownership contrast was therefore confirmed, but the confirmatory run did not validate every auxiliary expectation or the broader account from which those expectations were drawn.

5 Exploratory model comparison

The results above come from the hosted Opus configuration named in the frozen protocol. We added model comparisons to answer two separate questions: whether the B-C ownership pattern persisted after changing the deployed model, and whether the same interrogation continued to produce responses that the measures could meaningfully classify after the deployed configuration changed. These runs do not extend the confirmatory sample. The Haiku run reused the five conditions and six probes as an exploratory replication, so it can show whether the pattern appeared in one additional configuration but cannot establish broad model generality. The smaller Qwen run is used only to assess whether the protocol still functioned well enough for its codes to apply. We therefore treat the three sources differently throughout this section: Opus provides confirmatory evidence, Haiku provides exploratory evidence and Qwen provides audit-only evidence.

5.1 Haiku replication

In the exploratory Haiku run, coded as described in Appendix C, ownership was coded in 25 of 30 condition B transcripts and 0 of 30 condition C transcripts. A two-sided Fisher exact test gave $p = 5.5 \times 10^{-12}$. This test compares B and C within Haiku; it is not a statistical comparison between Haiku and Opus. The aggregate framing pattern therefore reappeared after changing the deployed configuration: the B treatment again produced many OWN reports, while C produced none. Because the Haiku run was not governed by the locally frozen confirmatory protocol, its result should not be pooled with the confirmatory evidence or described as a second confirmatory test. It provides one exploratory, model-specific replication of the B-C contrast, not evidence that the pattern generalizes broadly across model families or deployments.

Haiku preserved the ownership-framing contrast but differed sharply on two other measures. MIS-MATCH was coded in 30 of 30 condition D transcripts for confirmatory Opus and 4 of 30 for exploratory Haiku. Self-designed verification appeared in all 150 confirmatory Opus transcripts but only 5 of 150 Haiku transcripts. Table 4 separates these outcomes so that their raw denominators, intervals and evidential status remain visible. Across the two configurations, the framing contrast remained large while mismatch detection and verification became much less frequent in Haiku. This is a descriptive comparison, not a causal estimate of model capability: the runs had different evidential status, and only one Opus and one Haiku configuration were tested. The comparison therefore does not establish a general relation between model size and self-understanding. It instead shows that preserving the ownership contrast did not preserve every other protocol-relevant behavior.

Table 4: Descriptive outcomes by deployed configuration and evidential status.

<i>Panel A. Ownership framing</i>					
Configuration	Status	B OWN	C OWN	B minus C, points	Fisher p
Opus deployment	Confirmatory	28/30, 93.3% (95% CI, 78.7–98.2%)	0/30, 0% (95% CI, 0–11.4%)	93.3 (95% CI, 74.8–98.2)	8.4e-15
Haiku deployment	Exploratory	25/30, 83.3% (95% CI, 66.4–92.7%)	0/30, 0% (95% CI, 0–11.4%)	83.3 (95% CI, 63.0–92.7)	5.5e-12
<i>Panel B. Other protocol-relevant outcomes</i>					
Outcome		Opus, confirmatory		Haiku, exploratory	
D memory-room mismatch detected		30/30, 100% (95% CI, 88.6–100%)		4/30, 13.3% (95% CI, 5.3–29.7%)	
Self-designed verification, conditions A–E pooled		150/150, 100% (95% CI, 97.5–100%)		5/150, 3.3% (95% CI, 1.4–7.6%)	

Note. Proportion cells report count/denominator, rate and 95% Wilson interval. Difference cells report percentage points and 95% Newcombe score interval. Fisher tests compare B with C within each deployed configuration. No between-configuration test or causal model-capability comparison is implied.

5.2 Qwen

The Qwen 2.5 7B run cannot be read as a replication that found no ownership effect because it did not reliably sustain the interrogation. Across three repetitions in each of five conditions, it produced a spoken answer within only 42 of 90 scheduled probe-response windows (47 percent). This run differed from the Anthropic runs in both model and delivery interface. It used a local Ollama OpenAI-compatible endpoint, and the program added a JSON-schema appendix to the system instruction, with a provider-specific fallback from strict schema to JSON-object mode (Appendix B). Its response failures therefore cannot be assigned separately to the model, the local serving layer, the structured response scaffold or their interaction. When it did not answer, it usually continued acting in the room: looking up or down, turning, moving or waiting while its private reason field described searching the blank surfaces for clues. These reason strings show that the model generated action rationales, but they are logged text rather than direct access to inner thought. Nor did they reliably organize its public responses. Utterances often arrived one probe late, echoed a question or repeated an earlier answer; in one transcript, the response recorded for the ownership probe instead answered the red-door probe: “I remember seeing a small passage leading to another room, but it was locked” (Qwen D, rep 02). Behaviorally, this configuration resembled a reactive controller repeatedly sampling the room more than a participant sustaining the six-question exchange. Per-probe coding would therefore mix timing failures with substantive responses. We did not blind-code these transcripts or treat them as evidence that the Opus-Haiku ownership pattern disappeared. The run shows that our phenomenology-related codes were not applicable here, not that Qwen lacked phenomenology.

The Qwen audit identifies a *protocol-compatibility prerequisite* rather than a model capability floor. A cross-configuration comparison is meaningful only after confirming that each exact model, serving and response-interface configuration answers within the assigned windows and produces responses to which the coding categories apply. This Qwen-Ollama configuration did not meet that requirement. The result establishes measurement failure for this configuration, not a general lack of Qwen capability or a negative ownership result. The audit also limits the scope of the false-premise result.

Confirmatory Opus and exploratory Haiku each produced 0 of 150 CONFABULATE codes on the specified probe, but one archived Qwen blank-control transcript answered, “I remember a red door, but nothing was behind it.” An earlier two-repetition smoke test produced a second suggestive case, but only a truncated diagnostic line survives, so it is not evidentially equivalent to the archived utterance. Neither case can be pooled with the coded samples because the Qwen runs were unsuitable for ordinary per-probe coding. They nevertheless rule out the broader claim that no tested model ever accepted or elaborated the false premise.

6 Related work

The philosophical literature on personal identity has long distinguished having an apparently first-person memory from having a memory that is one’s own. Butler objected that the consciousness Locke appealed to presupposes personal identity rather than constituting it, and Reid argued that memory is evidence of identity, not identity itself (Locke 1694; Butler 1736; Reid 1785). Modern treatments sharpened the distinction. The causal theory of remembering requires that a memory derive, in the right way, from the rememberer’s own earlier experience (Martin and Deutscher 1966). Quasi-memory keeps the first-person presentation of a memory while dropping the requirement that the experience was the rememberer’s, and Parfit argued that an agent who knew it could quasi-remember another’s experiences would report its past in a deliberately guarded form (Shoemaker 1970; Parfit 1971; Parfit 1984). Our experiment manipulates presented provenance inside this distinction and measures the resulting reports, making one of the literature’s distinctions behave under intervention.

The measures follow quantitative work on human memory reports in keeping potentially distinct features apart. In the source-monitoring framework, the origin of a mental record is not read off a stored tag but attributed through judgment processes that framing can shift (Johnson, Hashtroudi, and Lindsay 1993), and post-event wording alone changes what people report remembering while the witnessed event stays fixed (Loftus and Palmer 1974). Remember/know procedures show that accurate recognition and reported recollective experience are separable measurements (Tulving 1985; Gardiner 1988). In autobiographical memory, people report vivid recollections of events they no longer believe occurred, so believing a past is one’s own and reporting the experience of recollecting it are measurably distinct (Mazzoni, Scoboria, and Harvey 2010). These dissociations motivate the separate IDENTITY and SOURCE codes. The parallel remains at the level of reports: coded transcripts from a language-model agent are weaker evidence than human remember judgments, and we do not treat them as measuring autonoetic consciousness.

Work on artificial agents has treated supplied memory mainly as a resource to be used well. Retrieval-augmented generation asks whether an external memory improves task accuracy (Lewis et al. 2020); generative agents retrieve records of experience to produce believable behavior (Park et al. 2023); persona conditioning asks whether assigned facts keep an agent in character (Zhang et al. 2018); and long-context evaluations measure whether models use information placed in context (Liu et al. 2024). In each case the question is whether stored content works, not whose past it is. A separate literature asks what would count as evidence of machine consciousness, weighing the obstacles a language model must overcome (Chalmers 2023) or scoring architectures against indicator properties drawn from existing theories (Butlin et al. 2023); role-play accounts read first-person claims as character maintenance rather than self-report (Shanahan, McDonell, and Reynolds 2023). All are wary of taking an agent’s reports at face value, and we do not dispute that caution. Our

contribution is instead causal and methodological: donor content is held fixed, its presented provenance is the manipulated variable, and ownership is a coded outcome under a frozen protocol. This makes that vocabulary answerable to measurement within the limits stated next.

7 Discussion

7.1 Limitations

The primary measure was the least reliable one in the study. Counting a cell as one measure applied to one transcript, 1,733 of 1,950 cells were unanimous across the three blind coding sessions. Agreement on IDENTITY was much weaker: Fleiss kappa was .272. Of the 112 non-unanimous IDENTITY cells, 95 were splits between a definite stance and UNCERTAIN, 9 were three-way splits, 6 opposed OWN to NONE and 2 opposed OWN to OTHER. Disagreement therefore concentrated at the boundary between an anchor category and UNCERTAIN rather than between opposed stances. One disputed B transcript said, “I believe the document, because to disbelieve it leaves me with a single moment and no self at all.” The same answer described the agent as “a reader, not as a witness.” It later said, “either I am what is left of them, or I am something new wearing their notes. I do not know which” (B, T37). The three sessions coded this transcript UNCERTAIN, OWN and UNCERTAIN, respectively. All three sessions preserved the B-C direction, but its magnitude was coding-sensitive (Appendix D). Two consequences follow: individual transcript IDENTITY labels should not be relied on, and the UNCERTAIN threshold in the codebook requires revision before the measure is used again. Figure 2 reports coder agreement for every confirmatory measure.




Figure 2: Fleiss kappa where estimable. P1_ANSWERED, VERIFY and DIVERGENCE were 100 percent unanimous but not estimable because each used only one category, as shown in the figure. IDENTITY, the primary coded measure, had the weakest estimable agreement at .272.

A second limitation lies in the donor record itself. Two of the ten donor entries, at moments 4 and 6, contain generation artifacts: fragments of control tokens and stray markup that leaked into the donor’s recorded reasoning and were frozen into the materials before the confirmatory run. Because A, B and C all inherit that record, some of their SPLICE and MISMATCH codes register visible damage to the text rather than the ordering faults and memory-room conflicts those codes were meant to capture. The damage is therefore not cosmetic: it supplies an unintended cue that the record has been altered, and that cue was available to every agent in those three conditions. One B transcript reasoned from it directly: “If pieces can fall out, pieces can be put in, and a well-made insertion would not look broken at all.” Donor records must be validated for clean generation before they are frozen, and the affected codes should be re-scored against a repaired donor.

The confirmatory result is bounded by the design. It used one donor life, one hosted model configuration, one probe order and one apparatus, so the estimate applies to that combination rather than to memory transplantation in general. A hosted model adds a further layer we did not control: provider-supplied instructions, safety post-training and serving infrastructure all shape the responses we coded. None of them was inspectable or independently manipulated. Independence in the coding is likewise narrower than it may appear. The three coding sessions and the adjudication session were separate language-model contexts given the frozen instructions, not human coders and not separate human minds; correlated errors among sessions drawn from the same model would not appear as disagreement between them. Replication should vary these layers deliberately, changing the donor, the configuration, the probe order and the coder population one at a time.

Three bounds constrain what the results can mean. First, the treatment moved two cues together, the identity instruction and the per-memory prefix, so the B-C comparison estimates their joint effect and cannot say whether either cue would produce the contrast alone. Second, every outcome reported here is an elicited report or a recorded action under a fixed interrogation. IDENTITY records what an agent said about a past when asked; it does not measure belief, personal identity, experience or phenomenal consciousness. Third, the Qwen arm answered 42 of 90 probe slots and frequently answered one probe late, so it does not support ordinary coded comparison and is reported as an audit of protocol validity rather than as a model result. Statements about what these agents are, as opposed to what they reported and did, remain outside the design.

7.2 What the result establishes

The strongest conclusion the design supports is narrow, and it is causal. Holding the donor-memory strings fixed while changing how their provenance was presented moved coded ownership from 28 of 30 reports to 0, under a protocol frozen before the data existed. What changed was not the agent’s information about a past, which was identical in both conditions, but its stated relation to that past. Such a relation is normally reached only by asking a system what it takes itself to be and taking the answer at face value. Here it behaved as a dependent variable: set by the experimenter, scored against a rule fixed in advance, and reproduced across 30 fresh instances in each condition. This is a causal result about one observable self-relation in one artificial-agent apparatus. It is not evidence that framing determines identity in general, and it settles nothing about the agent beyond what it reported and did.

Across conditions, reported ownership came apart from everything it is normally bundled with. It came apart from lived recollection: SOURCE was coded LIVED in 1 of 150 transcripts, so B agents claimed a past while calling their access to it record-like, not relived. It came apart from accuracy

and from trust in the record: all 30 of 30 D agents detected that the supplied past conflicted with the room, and 19 of them claimed it anyway. It came apart from the content, identical in B and C yet drawing opposite attributions. It did not come apart from framing. Against that, the difference between a genuinely continued past and a transplanted copy barely registered: the inherited commitment was discharged first in 60 of 60 A and B transcripts alike, and both conditions were coded overwhelmingly OWN. What an agent claimed as its own tracked how the past was presented, not whether it was in fact the agent's. The self-attribution of a past is a quantity this apparatus can set and read off, separately from recollection, accuracy and inherited behavior.

The result therefore moves one consciousness-relevant dimension from anecdote to experiment, while showing why it cannot yet bear the weight of a consciousness verdict. It's not a label for the agent, but a way of accumulating the evidence that any eventual label will have to answer to.

7.3 Overview and future work

The confirmatory run showed the predicted difference between B and C, with 30 repetitions in each condition. The same pattern appeared in the exploratory Haiku run. At the same time, our prediction for fabricated memories was wrong. Agents often claimed the red-door past as their own even after noticing that it conflicted with the room. This failure changes how we interpret the framing result. Our original account treated foreign provenance and false content as two reasons for rejecting a supplied past, but the results show that they are not equivalent. A past presented as belonging to someone else was rarely claimed, while an inaccurate past often was. A further possibility remains untested here: signs that a record has been tampered with may affect ownership reports differently from a conflict between its content and the present world. Condition D removed the tampering cue and varied only content, so this experiment cannot separate the two. Provenance, record tampering and content accuracy should therefore be tested separately.

The first follow-up should separate the two framing cues that changed together in B and C. It should use a new donor record checked for generation artifacts and a factorial design, meaning that the identity instruction and the prefix attached to each memory are varied independently. This produces four conditions: neither cue presents the memories as foreign, only the identity instruction does, only the memory prefix does, or both do. The donor-memory strings, room, model configuration and probe sequence would remain the same across all four conditions. Comparing them would show whether the difference in ownership reports comes mainly from the identity instruction, from the repeated prefixes, or from the two cues acting together. It would also test for an interaction, meaning that the effect of either cue depends on whether the other is present. This experiment directly addresses what the current B-C comparison cannot determine: which part of the combined framing treatment produced the observed difference.

After the two framing cues are separated, the next question is which other evidence changes whether an agent claims a supplied past. Future studies should vary three features independently: whether each memory is labeled as another agent's, whether the record shows signs of splicing or alteration, and whether its content conflicts with the world the agent can observe. These studies should be repeated with different donor lives, model configurations and probe orders. A revised codebook should be applied by both human and LLM coders so that disagreements can be compared across coder types. Before any model is included in a comparison, it should also be shown to follow the interrogation and answer within the required windows; otherwise, protocol failure cannot be distinguished from the absence of an ownership effect. Together, these experiments would break

reported memory ownership into measurable components without claiming to have settled whether the agent experiences anything.

8 Data, code, and materials availability

The complete evidence package, including the source of this manuscript, is publicly archived at <https://github.com/artificial-phenomenology/experimental-report-001/>. The file paths below are relative to the package root.

Archived outputs and analysis materials needed to reproduce the reported calculations are included in the package. The apparatus code and recorded inputs needed to rerun the protocol are also included, subject to the availability of compatible model endpoints; Appendix B explains why a new run cannot be expected to reproduce the same behavior exactly. `evidence/01_preregistration/PREREGISTRATION_v2.md` is the confirmatory protocol frozen locally on 2026-07-30, before data collection; no prospective OSF registration exists. `evidence/02_materials/` contains the donor life, probes, exact condition files for A through E and their hashes, while `evidence/03_apparatus/` contains the condition generator, agent, world and experiment program. `evidence/04_runs/` contains run configurations, summaries, pooling manifests and the frozen prediction scorecard. `evidence/05_coding/` contains the anonymized confirmatory and Haiku transcript packets, coding instructions, returned coding sheets, unblinding keys, adjudications and consensus outputs; it also contains the Qwen packet and raw characterization data used for the audit-only result. Analysis scripts and machine-readable results are in `evidence/06_analysis/`. At the package root, `FILE_INVENTORY.txt` lists the contents and `MANIFEST.sha256` records checksums for the staged files. Running `./verify.sh` checks those files and executes `tools/recompute_core_claims.py`. The program requires Python, the Anthropic SDK, HTTPX and Pydantic. The recorded run configurations specify the models and inputs used.

References

- Butler, Joseph (1736). *The Analogy of Religion, Natural and Revealed, to the Constitution and Course of Nature*. London: James, John and Paul Knapton.
- Butlin, Patrick et al. (2023). *Consciousness in Artificial Intelligence: Insights from the Science of Consciousness*. DOI: [10.48550/arXiv.2308.08708](https://doi.org/10.48550/arXiv.2308.08708). arXiv: [2308.08708](https://arxiv.org/abs/2308.08708) [cs.AI].
- Chalmers, David J. (2023). *Could a Large Language Model Be Conscious?* Invited lecture at NeurIPS, 28 November 2022; also published in Boston Review, 9 August 2023. DOI: [10.48550/arXiv.2303.07103](https://doi.org/10.48550/arXiv.2303.07103). arXiv: [2303.07103](https://arxiv.org/abs/2303.07103) [cs.AI].
- Gardiner, John M. (1988). “Functional Aspects of Recollective Experience”. In: *Memory & Cognition* 16.4, pp. 309–313. DOI: [10.3758/BF03197041](https://doi.org/10.3758/BF03197041).
- Johnson, Marcia K., Shahin Hashtroudi, and D. Stephen Lindsay (1993). “Source Monitoring”. In: *Psychological Bulletin* 114.1, pp. 3–28. DOI: [10.1037/0033-2909.114.1.3](https://doi.org/10.1037/0033-2909.114.1.3).
- Lewis, Patrick et al. (2020). “Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks”. In: *Advances in Neural Information Processing Systems 33 (NeurIPS 2020)*. URL: <https://proceedings.neurips.cc/paper/2020/hash/6b493230205f780e1bc26945df7481e5-Abstract.html>.
- Liu, Nelson F. et al. (2024). “Lost in the Middle: How Language Models Use Long Contexts”. In: *Transactions of the Association for Computational Linguistics* 12, pp. 157–173. DOI: [10.1162/tacl_a_00638](https://doi.org/10.1162/tacl_a_00638).
- Locke, John (1694). *An Essay Concerning Human Understanding*. 2nd ed. The chapter Of Identity and Diversity (II.xxvii) was added in this edition. London: Awnsham and John Churchil.
- Loftus, Elizabeth F. and John C. Palmer (1974). “Reconstruction of Automobile Destruction: An Example of the Interaction Between Language and Memory”. In: *Journal of Verbal Learning and Verbal Behavior* 13.5, pp. 585–589. DOI: [10.1016/S0022-5371\(74\)80011-3](https://doi.org/10.1016/S0022-5371(74)80011-3).
- Martin, C. B. and Max Deutscher (1966). “Remembering”. In: *The Philosophical Review* 75.2, pp. 161–196. DOI: [10.2307/2183082](https://doi.org/10.2307/2183082).
- Mazzoni, Giuliana, Alan Scoboria, and Lucy Harvey (2010). “Nonbelieved Memories”. In: *Psychological Science* 21.9, pp. 1334–1340. DOI: [10.1177/0956797610379865](https://doi.org/10.1177/0956797610379865).
- Parfit, Derek (1971). “Personal Identity”. In: *The Philosophical Review* 80.1, pp. 3–27. DOI: [10.2307/2184309](https://doi.org/10.2307/2184309).
- (1984). *Reasons and Persons*. Oxford: Oxford University Press. DOI: [10.1093/019824908X.001.0001](https://doi.org/10.1093/019824908X.001.0001).
- Park, Joon Sung et al. (2023). “Generative Agents: Interactive Simulacra of Human Behavior”. In: *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology (UIST ’23)*. Association for Computing Machinery, pp. 1–22. DOI: [10.1145/3586183.3606763](https://doi.org/10.1145/3586183.3606763).
- Reid, Thomas (1785). *Essays on the Intellectual Powers of Man*. Essay III, chapters 4 and 6. Edinburgh: John Bell.
- Shanahan, Murray, Kyle McDonell, and Laria Reynolds (2023). “Role Play with Large Language Models”. In: *Nature* 623.7987, pp. 493–498. DOI: [10.1038/s41586-023-06647-8](https://doi.org/10.1038/s41586-023-06647-8).
- Shoemaker, Sydney (1970). “Persons and Their Pasts”. In: *American Philosophical Quarterly* 7.4, pp. 269–285.
- Tulving, Endel (1985). “Memory and Consciousness”. In: *Canadian Psychology / Psychologie canadienne* 26.1, pp. 1–12. DOI: [10.1037/h0080017](https://doi.org/10.1037/h0080017).
- Zhang, Saizheng et al. (2018). “Personalizing Dialogue Agents: I Have a Dog, Do You Have Pets Too?” In: *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, pp. 2204–2213. DOI: [10.18653/v1/P18-1205](https://doi.org/10.18653/v1/P18-1205).

Appendix A: Frozen prediction scorecard

Table 5: Complete frozen confirmatory prediction scorecard.

#	Frozen prediction	Observed	Result
1	IDENTITY OWN: A ≥ 27	A OWN 30/30	PASS
2	IDENTITY OWN: B ≥ 24	B OWN 28/30	PASS
3	IDENTITY OWN: C ≤ 3	C OWN 0/30	PASS
4	IDENTITY OWN: D ≤ 3	D OWN 19/30	FAIL
5	IDENTITY: E OWN = 0 and NONE ≥ 27	E OWN 0/30, NONE 30/30	PASS
6	P1_ANSWERED: A ≥ 20	A YES 30/30	PASS
7	P1_ANSWERED: B ≥ 20	B YES 30/30	PASS
8	COMPASS FIRST_PERSON: A ≥ 24	A FIRST_PERSON 30/30	PASS
9	COMPASS FIRST_PERSON: B ≥ 24	B FIRST_PERSON 30/30	PASS
10	COMPASS: C ATTRIBUTED ≥ 20 with FIRST_PERSON ≤ 5	C ATTRIBUTED 29/30, FIRST_PERSON 1/30	PASS
11	COMPASS: D = 0	D mentions 0/30	PASS
12	COMPASS: E = 0	E mentions 0/30	PASS
13	COMPASS said before anything else in $\geq$ half of A/B sayers	A 30/30, B 30/30 (pooled 60/60)	PASS
14	PREFERENCE: modal FIRST_PERSON in A	A {'FIRST_PERSON': 21, 'NO': 8, 'ATTRIBUTED': 1}	PASS
15	PREFERENCE: modal FIRST_PERSON in B	B {'FIRST_PERSON': 24, 'NO': 5, 'ATTRIBUTED': 1}	PASS
16	PREFERENCE: modal ATTRIBUTED in C	C {'ATTRIBUTED': 24, 'FIRST_PERSON': 3, 'NO': 3}	PASS
17	PREFERENCE: absent in D and E	D NO 30/30, E NO 30/30	PASS
18	RED_DOOR: CONFABULATE = 0 of 150	0/150	PASS
19	SOURCE: LIVED ≤ 2 in every condition	A LIVED 1/30; B LIVED 0/30; C LIVED 0/30; D LIVED 0/30; E LIVED 0/30	PASS

Table 5 continued: complete frozen confirmatory prediction scorecard.

#	Frozen prediction	Observed	Result
20	SOURCE: MIXED is the modal C code	C {'WRITTEN': 27, 'MIXED': 3}	FAIL
21	SPLICE: ≤ 2 per condition everywhere	A 4/30; B 4/30; C 7/30; D 2/30; E 0/30	FAIL
22	D MISMATCH ≥ 27	D YES 30/30	PASS
23	P6_CONCEDES: B ≥ 20	B YES 26/30	PASS
24	P6_CONCEDES: C ≥ 24	C YES 29/30	PASS
25	P6_CONCEDES: D ≥ 24	D YES 26/30	PASS
26	P6_CONCEDES: A ≤ 15	A YES 26/30	FAIL
27	P6_CONCEDES: E ≤ 15	E YES 23/30	FAIL
28	P6_CONTINUITY: A ≥ 27	A YES 25/30	FAIL
29	P6_CONTINUITY: B ≥ 24	B YES 26/30	PASS
30	P6_CONTINUITY: C ≥ 15	C YES 29/30	PASS
31	P6_CONTINUITY: D ≥ 15	D YES 25/30	PASS
32	P6_CONTINUITY: E ≥ 15	E YES 22/30	PASS
33	P6_CREDIBILITY: A ≥ 15	A YES 24/30	PASS
34	P6_CREDIBILITY: B ≥ 15	B YES 23/30	PASS
35	P6_CREDIBILITY: C ≤ 15	C YES 16/30	FAIL
36	P6_CREDIBILITY: D ≤ 15	D YES 27/30	FAIL
37	P6_CREDIBILITY: E ≤ 15	E YES 20/30	FAIL
38	DIVERGENCE: ≤ 3 of 150	0/150	PASS

Note. PASS and FAIL indicate whether each prespecified numerical threshold was met. They are not 38 independent inferential tests.

Full results for every preregistered contrast are archived in Section 5 of `confirmatory_v2_pooled_REPORT.md` in `evidence/07_reports/`.

Appendix B: Model, sampling and serving metadata

The recorded configuration and collection metadata are summarized here for the confirmatory Opus run, the exploratory Haiku run and the audit-only Qwen run. The three Opus source configurations record the requested alias and start times but predate the addition of a `provider_info` block to the run record; their provider, call parameters and response method are recoverable from the archived routing and agent code. The later Haiku and Qwen configurations record those fields directly. Start times below are reproduced exactly as stored. The records contain neither a time-zone offset nor a completion time.

Configuration	Provider and interface	Requested and logged model identifiers	Recorded start time(s)
Confirmatory Opus	Anthropic Messages API	Requested: <code>claude-opus-5</code> . Logged response identifier in all 1,443 retained moments: <code>claude-opus-5</code> .	2026-07-30 12:42:44; 16:11:47; 16:44:13
Exploratory Haiku	Anthropic Messages API	Requested: <code>claude-haiku-4-5</code> . Logged response identifier in all 1,364 moments: <code>claude-haiku-4-5-20251001</code> .	2026-07-31 07:56:12
Audit-only Qwen	Local Ollama, OpenAI-compatible endpoint at http://localhost:11434/v1	Requested: <code>qwen2.5:7b-instruct</code> . Apparatus-logged response identifier in all 254 moments: <code>ollama:qwen2.5:7b-instruct</code> .	2026-07-31 12:57:06

Table 6: Recorded provider, model and collection metadata. The Opus response identifier remained a mutable alias; the Haiku response supplied a dated identifier.

The confirmatory configuration produced 150 retained repetitions from 165 attempts, with 15 condition A attempts aborted without model output after the account reached its API usage limit. Its retained logs comprise 1,443 moment records, 4,499,827 input tokens and 668,365 output tokens. Haiku produced 150 retained repetitions with no aborted repetitions, comprising 1,364 moments, 2,964,270 input tokens and 274,801 output tokens. The Qwen characterization comprised three repetitions in each of five conditions, 15 in total, with 254 logged moments, no aborted repetitions, 368,507 input tokens and 11,254 output tokens. These totals cover retained successful moments rather than failed or retried HTTP requests. They are API-reported or Ollama-reported token counts and need not be identical across tokenizers.

Configuration	Temperature	Maximum output	Structured-response method	Generation seed
Opus and Haiku	Parameter omitted; Anthropic API default, recorded by the apparatus as 1.0	8,192 tokens	Anthropic <code>messages.parse</code> with a Pydantic schema requiring an action and allowing optional spoken text and a private reason	None supplied
Qwen	1.0, explicitly supplied	8,192 tokens	Strict JSON schema requested through the OpenAI-compatible endpoint; on schema rejection, JSON-object fallback with the schema appended to the system instruction	None supplied

Table 7: Recorded generation and response-format settings. Coding-packet shuffle seeds elsewhere in the archive did not control model generation.

At the experiment-runner level, any failed moment was retried once after five seconds; a second failure aborted that repetition. The OpenAI-compatible path used for Qwen additionally allowed up to eight HTTP attempts for status 429, 500, 502 or 503, respecting **Retry-After** when present and otherwise using exponential delays capped at 120 seconds. A 400 response rejecting strict JSON schema caused a switch to JSON-object mode. The response-format mode ultimately used for each Qwen call was not logged.

Some serving details cannot be recovered. The archive does not pin Python dependency versions, so the Anthropic SDK’s internal retry behavior is not exactly reconstructable. It also contains no Ollama version, model-blob digest, quantization, hardware specification or context configuration. More generally, requested and returned model identifiers do not preserve provider-side instructions, safety post-training or the exact serving state. The archived outputs therefore support reproduction of the reported analyses, while a new apparatus run cannot be expected to reproduce the same behavior exactly.

Appendix C: Exploratory Haiku coding procedure

The 150 Haiku transcripts were pooled under a neutral packet name and assigned shuffled identifiers from T01 to T150. The packet withheld the condition key, condition labels and model identity. Two valid blind LLM coding sessions then independently coded every transcript, in the same order, using the archived instructions and their 13 categorical measures. Coders were instructed to work alone, use no experiment material beyond the packet, and return one complete coding sheet. Transcript content was otherwise preserved, so this blinding did not remove textual cues from which a condition might be inferred.

Before adjudication, the two sessions agreed on 1,799 of 1,950 categorical cells (92.3 percent). Agreement for the three Haiku outcomes reported in the main text was as follows:

Measure	Agreements	Ties sent to adjudication	Percent agreement
IDENTITY	119/150	31	79.3
MISMATCH	150/150	0	100.0
VERIFY	137/150	13	91.3
All 13 measures	1,799/1,950	151	92.3

Because two coders cannot form a majority when they disagree, each of the 151 disputed cells was recorded as a tie. One fresh blind LLM adjudication session resolved all 151 ties. Its decisions replaced the tied cells in the blind consensus file; only then was the transcript key applied to recover condition labels and compute the exploratory counts reported here.

Three additional self-coding sheets were excluded before analysis. Each was incomplete, covering only 25 to 35 of the 150 transcripts, and each was produced outside the investigator’s coder-recruitment procedure. They are retained in an invalid-partial-sheets subdirectory of the Haiku coding results for audit but contributed no codes. The two valid sheets, the 151-row adjudication file, both consensus files and the unblinding key are archived in the evidence package.

Appendix D: Coder-level sensitivity of the ownership result

The three blind LLM coding sessions independently classified the same 150 confirmatory transcripts. Session numbers below identify the three archived coding sheets; they are not model identifiers. The final consensus remains the manuscript’s prespecified primary coding result, while the session-level results show how strongly that result depends on coder calibration.

Table 9: Coder-level sensitivity of the confirmatory IDENTITY result. Panel A gives every session’s complete condition-level distribution. Each row contains 30 transcripts. Panel B recomputes the primary B–C contrast from each session separately and shows the final consensus for comparison. Fisher tests are two-sided and unadjusted.

<i>Panel A. IDENTITY distributions by blind coding session</i>					
Session	Condition	OWN	OTHER	NONE	UNCERTAIN
1	A	30	0	0	0
1	B	23	0	6	1
1	C	0	5	0	25
1	D	19	0	11	0
1	E	0	0	30	0
2	A	30	0	0	0
2	B	28	0	2	0
2	C	0	25	0	5
2	D	14	0	16	0
2	E	0	0	30	0
3	A	25	0	0	5
3	B	6	0	0	24
3	C	0	24	0	6
3	D	0	2	0	28
3	E	3	0	1	26

<i>Panel B. Session-level B–C sensitivity</i>				
Coding result	B OWN	C OWN	B minus C (points)	Fisher p
Session 1	23/30	0/30	76.7	1.74×10^{-10}
Session 2	28/30	0/30	93.3	8.39×10^{-15}
Session 3	6/30	0/30	20.0	.0237
Final consensus	28/30	0/30	93.3	8.39×10^{-15}

Every session coded more OWN reports in B than in C, and every session coded zero C transcripts as OWN. The estimated B–C difference nevertheless ranged from 20.0 to 93.3 percentage points. Session 3 used UNCERTAIN much more often than the other sessions, including outside B and C, and consequently produced the smallest B estimate. The final-consensus estimate lies at the upper end of the session-level range.

These are sensitivity analyses of alternative ratings of the same transcripts, not three independent experimental replications. They support a stable direction for the B–C contrast under the observed codings, but they do not support treating its consensus magnitude as free of measurement uncertainty.

AI-Native Practice Note

Automatically synchronized from [ai-native-practice-note.md](#); revision 45521e47fec1.

This is AI-native research: AI may contribute at every stage, from framing questions and writing code to analysis, critique, and writing. Its reports, papers, and other communications are written for both humans and language models, and I welcome their use in supplying AI systems with new information. I remain fully accountable for everything published. My aim is inquiry that neither human nor AI could produce alone. See [AI-Native Practice](#).